

Nebula GPU Interconnect System: Technical Architecture and Implementation

Pranav Chandra, Pramit Pal, Meghadri Ghosh
Team Bob

September 2025

Project Overview

- **Goal:** Scalable, cache-coherent GPU interconnect for AI/ML workloads
- **Topology:** 2D mesh, 2x2 to 8x8 grid (up to 64 GPUs)
- **Protocols:** ARM AMBA AXI4 (non-coherent), CHI (coherent)
- **Implementation:** SystemVerilog RTL, Python analysis, Web dashboard
- **Key Features:**
 - Five-stage router pipeline with virtual channels
 - Adaptive and deterministic routing
 - Protocol translation bridges (AXI4/CHI \leftrightarrow NoC)
 - Performance monitoring and visualization

System Architecture

- Mesh topology generated by `nebula_mesh_top.sv`
- Routers interconnected, edge nodes terminated
- Each node: router + protocol bridge + GPU/memory interface

Router Microarchitecture

Five-Stage Pipeline:

- 1 Buffer Write (BW): Input port, VC allocation
- 2 Route Computation (RC): XY/adaptive routing
- 3 Virtual Channel Allocation (VA): VC state machines, credit reservation
- 4 Switch Allocation (SA): Crossbar arbitration, fairness
- 5 Switch Traversal (ST): Data transmission, credit management

Features:

- 4 VCs per port, 16-flit FIFO depth
- Credit-based flow control
- Round-robin arbitration
- Congestion-aware adaptive routing

Routing Algorithms

- **XY Routing:** Dimension-ordered, deadlock-free
- **Adaptive Routing:** Congestion-aware, selects least congested path
- **Implementation:**
 - Combinational logic compares router and destination coordinates
 - Congestion metrics: buffer utilization, credit counts
 - Programmable thresholds for adaptive switching

Technical Details

- Route Computation: lines 260-380 in `nebula_router.sv`
- VC Allocation: lines 630-690 in `nebula_router.sv`
- Switch Allocation: lines 460-570 in `nebula_router.sv`

Flit and Packet Format

- **Flit:** 256 bits, defined in `nebula_pkg.sv`
- **Header Fields:**
 - Type (HEAD/BODY/TAIL/SINGLE)
 - Source/Dest coordinates (4 bits each)
 - VC ID (2 bits)
 - Sequence number (16 bits)
 - Packet ID (8 bits)
 - QoS (4 bits)
 - CRC (32 bits)
- **Payload:** 208 bits
- **Multi-flit packets:** HEAD, BODY, TAIL; single-flit packets supported

Protocol Translation Bridges

- **AXI-NoC Bridge** (`nebula_axi_noc_bridge.sv`):
 - Burst decomposition, address mapping
 - 64-entry reorder buffer
 - Packet assembly/disassembly
 - Error detection, latency monitoring
- **CHI-NoC Bridge** (`nebula_chi_noc_bridge.sv`):
 - CHI message classification, VC mapping
 - Directory-based MOESI coherency
 - Snoop response aggregation
 - Timeout/error handling
- **Outstanding Transaction Management:**
 - Hardware table tracks up to 64 operations
 - Transaction ID, address, burst length, sequence, timeout

Packet Assembly and Disassembly

- **Assembler** (`nebula_packet_assembler.sv`):
 - Converts protocol transactions to flit packets
 - Header generation, payload segmentation
 - Address-to-coordinate mapping
- **Disassembler** (`nebula_packet_disassembler.sv`):
 - Reconstructs transactions from flits
 - CRC verification, sequence management
 - Handles out-of-order and multi-path delivery

Credit-Based Flow Control

- **Credit Controller** (`nebula_credit_flow_ctrl.sv`):
 - Per-VC credit counters, max 16
 - Increment on flit acceptance, decrement on allocation
 - Prevents buffer overflow, deadlock
 - `credits_available` signals for arbitration
- **Flow Control Protocol:**
 - Sender stalls if `credits = 0`
 - Lossless operation guaranteed

System Integration

- **Top-Level Modules:**

- `nebula_system_top.sv`: Mesh instantiation, AXI4 interfaces, address mapping
- `nebula_mesh_top.sv`: Router grid generation, edge handling
- `nebula_top.sv`: Configuration registers, system status, debug

- **Address Mapping:**

- 64-bit global address space
- Static/dynamic schemes (partitioning, hashing)
- Hardware decoder for coordinate extraction

- **Collective Operations:**

- Broadcast/multicast via packet replication
- Hardware barriers, reduction support

Verification and Testing

- **RTL Testbenches** (`code/tb/`):

- Unit tests: packet assembler/disassembler, FIFO, arbiter
- Integration tests: AXI-NoC bridge, router, mesh
- System-level: end-to-end, multi-hop, congestion, performance

- **Python Tools:**

- `nebula_traffic_generator.py`: Traffic pattern generation, testbench synthesis
- `nebula_vcd_parser.py`: VCD trace analysis, packet event extraction

- **Metrics:**

- Latency, throughput, congestion, error rates
- Performance counters in routers and system top

Web Dashboard

- **Backend:** Flask + Socket.IO (web_dashboard/backend/app.py)
- **Frontend:** Vanilla JS + Vite + Tailwind (web_dashboard/frontend/)
- **Features:**
 - Real-time mesh visualization (SVG)
 - Animated packet flows, congestion heatmap
 - Performance metric graphing (utilization, latency, throughput)
 - VCD trace replay, simulation control
 - Traffic pattern selection (Uniform, Hotspot, CNN, Matrix, Transformer, etc.)
- **Integration:**
 - Runs Verilog simulations, parses VCD, updates UI via WebSocket
 - API endpoints for mesh, performance, simulation control

Performance Monitoring

- **Router-Level Counters:**

- Packets forwarded per direction
- Buffer utilization, VC statistics
- Congestion, temperature (simulated)

- **System-Level Metrics:**

- Total packets routed, average latency, throughput
- Error and protocol violation tracking
- Historical trending for optimization

- **Visualization:**

- Time-series graphs, mesh overlays
- VCD replay for cycle-accurate analysis

Technical Challenges

- **Scalability:** Mesh generation, address mapping for large grids
- **Coherency:** CHI protocol edge cases, directory state management
- **Adaptive Routing:** Congestion metrics, deadlock avoidance
- **Verification:** Multi-level test coverage, error injection
- **Integration:** Protocol bridge correctness, performance counter aggregation
- **Visualization:** Real-time updates, VCD parsing, UI responsiveness

Current Status and Future Work

- **Completed:**

- Mesh topology, router pipeline, protocol bridges
- AXI4/CHI support, packet assembler/disassembler
- Python traffic generator, VCD parser, web dashboard
- Verification infrastructure, performance monitoring

- **In Progress:**

- Advanced adaptive routing, QoS features
- Hierarchical clustering, multi-mesh support
- Enhanced CHI edge case handling
- Dashboard UI improvements, analytics

- **Planned:**

- FPGA prototyping, hardware deployment
- Deep learning workload benchmarks
- Fault tolerance, dynamic reconfiguration

References

- SystemVerilog RTL: `code/rtl/`
- Python Tools: `code/python/`
- Testbenches: `code/tb/`
- Documentation: `docs/final_report.tex`, `docs/abstract.tex`
- Dashboard: `web_dashboard/`

Questions?