

Nebula GPU Interconnect System: Technical Architecture and Implementation

Pranav Chandra, Pramit Pal, Meghadri Ghosh
Team Bob

September 2025

Project Overview

- **Goal:** Scalable, cache-coherent GPU interconnect for AI/ML workloads
- **Topology:** 2D mesh, 2x2 to 8x8 grid (up to 64 GPUs)
- **Protocols:** ARM AMBA AXI4 (non-coherent), CHI (coherent)
- **Implementation:** SystemVerilog RTL, Python analysis, Web dashboard
- **Key Features:**
 - Five-stage router pipeline with virtual channels
 - Adaptive and deterministic routing
 - Protocol translation bridges (AXI4/CHI \leftrightarrow NoC)
 - Performance monitoring and visualization

System Architecture

- Mesh topology generated by `nebula_mesh_top.sv`
- Routers interconnected, edge nodes terminated
- Each node: router + protocol bridge + GPU/memory interface




Figure: Nebula Mesh Topology Example (4x4)

Node Connections

- Nodes connect to their N, S, E & W neighbours via bidirectional links.
- The links carry flit data, control signals, and credit information.
- 2D mesh architecture, scalable from 2x2 mesh to 8x8 mesh.



Figure: Mesh Topology Example from Web Dashboard (8x8)

Router Microarchitecture

Five-Stage Pipeline:

- 1 Buffer Write (BW): Input port, VC allocation
- 2 Route Computation (RC): XY/adaptive routing
- 3 Virtual Channel Allocation (VA): VC state machines, credit reservation
- 4 Switch Allocation (SA): Crossbar arbitration, fairness
- 5 Switch Traversal (ST): Data transmission, credit management

Features:

- 4 VCs per port, 16-flit FIFO depth
- Credit-based flow control
- Round-robin arbitration
- Congestion-aware adaptive routing

Routing Algorithms

- **XY Routing:** Dimension-ordered, deadlock-free
- **Adaptive Routing:** Congestion-aware, selects least congested path
- **Implementation:**
 - Combinational logic compares router and destination coordinates
 - Congestion metrics: buffer utilization, credit counts
 - Programmable thresholds for adaptive switching

Technical Details

- Route Computation: lines 260-380 in `nebula_router.sv`
- VC Allocation: lines 630-690 in `nebula_router.sv`
- Switch Allocation: lines 460-570 in `nebula_router.sv`

Flit and Packet Format

- **Flit:** 256 bits, defined in `nebula_pkg.sv`
- **Header Fields:**
 - Type (HEAD/BODY/TAIL/SINGLE)
 - Source/Dest coordinates (4 bits each)
 - VC ID (2 bits)
 - Sequence number (16 bits)
 - Packet ID (8 bits)
 - QoS (4 bits)
 - CRC (32 bits)
- **Payload:** 208 bits
- **Multi-flit packets:** HEAD, BODY, TAIL; single-flit packets supported

Protocol Translation Bridges

- **AXI-NoC Bridge** (`nebula_axi_noc_bridge.sv`):
 - Burst decomposition, address mapping
 - 64-entry reorder buffer
 - Packet assembly/disassembly
 - Error detection, latency monitoring
- **CHI-NoC Bridge** (`nebula_chi_noc_bridge.sv`):
 - CHI message classification, VC mapping
 - Directory-based MOESI coherency
 - Snoop response aggregation
 - Timeout/error handling
- **Outstanding Transaction Management:**
 - Hardware table tracks up to 64 operations
 - Transaction ID, address, burst length, sequence, timeout

Packet Assembly and Disassembly

- **Assembler** (`nebula_packet_assembler.sv`):
 - Converts protocol transactions to flit packets
 - Header generation, payload segmentation
 - Address-to-coordinate mapping
- **Disassembler** (`nebula_packet_disassembler.sv`):
 - Reconstructs transactions from flits
 - CRC verification, sequence management
 - Handles out-of-order and multi-path delivery

Credit-Based Flow Control

- **Credit Controller** (`nebula_credit_flow_ctrl.sv`):
 - Per-VC credit counters, max 16
 - Increment on flit acceptance, decrement on allocation
 - Prevents buffer overflow, deadlock
 - `credits_available` signals for arbitration
- **Flow Control Protocol:**
 - Sender stalls if `credits = 0`
 - Lossless operation guaranteed

System Integration

- **Top-Level Modules:**

- `nebula_system_top.sv`: Mesh instantiation, AXI4 interfaces, address mapping
- `nebula_mesh_top.sv`: Router grid generation, edge handling
- `nebula_top.sv`: Configuration registers, system status, debug

- **Address Mapping:**

- 64-bit global address space
- Static/dynamic schemes (partitioning, hashing)
- Hardware decoder for coordinate extraction

- **Collective Operations:**

- Broadcast/multicast via packet replication
- Hardware barriers, reduction support

Verification and Testing

- **RTL Testbenches** (`code/tb/`):

- Unit tests: packet assembler/disassembler, FIFO, arbiter
- Integration tests: AXI-NoC bridge, router, mesh
- System-level: end-to-end, multi-hop, congestion, performance

- **Python Tools:**

- `nebula_traffic_generator.py`: Traffic pattern generation, testbench synthesis
- `nebula_vcd_parser.py`: VCD trace analysis, packet event extraction

- **Metrics:**

- Latency, throughput, congestion, error rates
- Performance counters in routers and system top

Web Dashboard

- **Backend:** Flask + Socket.IO (web_dashboard/backend/app.py)
- **Frontend:** Vanilla JS + Vite + Tailwind (web_dashboard/frontend/)
- **Features:**
 - Real-time mesh visualization (SVG)
 - Animated packet flows, congestion heatmap
 - Performance metric graphing (utilization, latency, throughput)
 - VCD trace replay, simulation control
 - Traffic pattern selection (Uniform, Hotspot, CNN, Matrix, Transformer, etc.)
- **Integration:**
 - Runs Verilog simulations, parses VCD, updates UI via WebSocket
 - API endpoints for mesh, performance, simulation control

Performance Monitoring

- **Router-Level Counters:**

- Packets forwarded per direction
- Buffer utilization, VC statistics
- Congestion, temperature (simulated)

- **System-Level Metrics:**

- Total packets routed, average latency, throughput
- Error and protocol violation tracking
- Historical trending for optimization

- **Visualization:**

- Time-series graphs, mesh overlays
- VCD replay for cycle-accurate analysis

Technical Challenges

- **Scalability:** Mesh generation, address mapping for large grids
- **Coherency:** CHI protocol edge cases, directory state management
- **Adaptive Routing:** Congestion metrics, deadlock avoidance
- **Verification:** Multi-level test coverage, error injection
- **Integration:** Protocol bridge correctness, performance counter aggregation
- **Visualization:** Real-time updates, VCD parsing, UI responsiveness

Current Status and Future Work

- **Completed:**

- Mesh topology, router pipeline, protocol bridges
- AXI4/CHI support, packet assembler/disassembler
- Python traffic generator, VCD parser, web dashboard
- Verification infrastructure, performance monitoring

- **In Progress:**

- Advanced adaptive routing, QoS features
- Hierarchical clustering, multi-mesh support
- Enhanced CHI edge case handling
- Dashboard UI improvements, analytics

- **Planned:**

- FPGA prototyping, hardware deployment
- Deep learning workload benchmarks
- Fault tolerance, dynamic reconfiguration

References

- SystemVerilog RTL: `code/rtl/`
- Python Tools: `code/python/`
- Testbenches: `code/tb/`
- Documentation: `docs/final_report.tex`, `docs/abstract.tex`
- Dashboard: `web_dashboard/`

Questions?

Nebula RTL Modules: Overview

- The Nebula NoC is composed of modular, parameterized SystemVerilog RTL blocks.
- Each module is designed for scalability, protocol flexibility, and high throughput.
- This section details the design decisions, code structure, and unique features of each RTL module.

nebula_top.sv: System Integration (1/2)

Design Decisions:

- Parameterized mesh size, protocol enable/disable, and config bus widths.
- Unified configuration and status interface for software control.
- Modular instantiation of routers, bridges, and monitoring logic.

Key Code:

- Parameter and function definitions for mesh configuration.
- Synchronous logic for system control and status.

nebula_top.sv: System Integration (2/2)

Cool Features:

- System-wide error aggregation and debug trace export.
- Performance counters and status registers for real-time monitoring.
- Generate blocks for scalable instantiation.

Key Code:

- Generate loops for node instantiation.
- Status aggregation logic.

nebula_system_top.sv: System-Level Integration (1/2)

Design Decisions:

- Centralizes mesh, memory, and protocol bridge instantiation.
- Implements global address decoding and mapping.
- Coordinates system resets and status propagation.

Key Code:

- Mesh instantiation with parameterized width and height.
- Address mapping logic from global to local address space.

nebula_system_top.sv: System-Level Integration (2/2)

Cool Features:

- Flexible address mapping for heterogeneous memory systems.
- System-level reset and error handling.
- Integration hooks for simulation and testbench control.

Key Code:

- Synchronous reset and control logic.
- Error status propagation and handling.

nebula_mesh_top.sv: Mesh Topology (1/2)

Design Decisions:

- 2D grid instantiation using nested generate loops.
- Edge node handling (tie-off or wrap-around).
- Parameterized for mesh width/height.

Key Code:

- Generate loops for X and Y dimensions of the mesh.
- Conditional logic for edge node connections.

nebula_mesh_top.sv: Mesh Topology (2/2)

Cool Features:

- Automated port wiring for all mesh neighbors.
- Edge detection logic for boundary nodes.
- Supports mesh reconfiguration for fault tolerance.

Key Code:

- Port connection logic using generate and if-else constructs.
- Tie-off or wrap-around logic for edge cases.

nebula_router.sv: Router Pipeline (1/2)

Design Decisions:

- Five-stage pipeline for high throughput.
- Virtual channel allocation and credit-based flow control.
- Adaptive and deterministic routing support.

Key Code:

- Synchronous logic for each stage of the router pipeline.
- Combinational logic for route computation and switch allocation.

nebula_router.sv: Router Pipeline (2/2)

Cool Features:

- Per-port and per-VC performance counters.
- Error detection for buffer overflow and protocol violations.
- Round-robin arbitration for switch allocation.

Key Code:

- Performance counter increment and error flag logic.
- Arbitration logic for switch allocation.

nebula_axi_if.sv: AXI4 Interface (1/2)

Design Decisions:

- Parameterized AXI widths for flexible integration.
- FIFO-based buffering for all AXI channels.
- Clock domain crossing support.

Key Code:

- AXI signal assignments and FIFO push/pop logic.
- Clock domain crossing synchronizers.

nebula_axi_if.sv: AXI4 Interface (2/2)

Cool Features:

- AXI burst handling and response reordering.
- Backpressure propagation to the NoC.
- Error injection for verification.

Key Code:

- Logic for burst-to-packet conversion and response assembly.
- Error detection and handling logic.

nebula_axi_noc_bridge.sv: AXI4-NoC Bridge (1/2)

Design Decisions:

- Decomposes AXI bursts into NoC packets.
- Maintains reorder buffer for out-of-order responses.
- Tracks outstanding transactions by ID.

Key Code:

- Logic for AXI to NoC packet conversion.
- Reorder buffer management for response assembly.

nebula_axi_noc_bridge.sv: AXI4-NoC Bridge (2/2)

Cool Features:

- Out-of-order response assembly from NoC flits.
- Error detection and reporting for protocol violations.
- Latency monitoring for each transaction.

Key Code:

- Logic for assembling responses from NoC flits.
- Error and latency monitoring logic.

nebula_chi_interface.sv: CHI Protocol (1/2)

Design Decisions:

- Implements CHI request, response, and data channels.
- State machine for protocol sequencing.
- Snoop and directory support.

Key Code:

- CHI protocol state machine implementation.
- Logic for handling CHI requests and responses.

nebula_chi_interface.sv: CHI Protocol (2/2)

Cool Features:

- FIFO-based buffering for all CHI channels.
- Protocol error detection and recovery.
- Integration with directory and snoop logic.

Key Code:

- Buffer management logic for CHI channels.
- Error handling and recovery logic.

nebula_chi_noc_bridge.sv: CHI-NoC Bridge (1/2)

Design Decisions:

- Maps CHI message types to NoC VCs.
- Handles coherency, snoop, and directory state transitions.
- Aggregates snoop responses and manages timeouts.

Key Code:

- Logic for mapping CHI messages to NoC virtual channels.
- Snoop response aggregation and timeout handling logic.

nebula_chi_noc_bridge.sv: CHI-NoC Bridge (2/2)

Cool Features:

- Timeout detection for snoop responses.
- Priority mapping for CHI traffic classes.
- Error handling for protocol mismatches.

Key Code:

- Timeout detection and handling logic.
- Priority encoding logic for CHI messages.

nebula_chi_directory.sv: Directory Controller (1/2)

Design Decisions:

- Implements MOESI directory for cache coherency.
- Handles snoop requests and state transitions.
- Manages race conditions and protocol corner cases.

Key Code:

- Directory state machine implementation for MOESI.
- Logic for handling snoop requests and responses.

nebula_chi_directory.sv: Directory Controller (2/2)

Cool Features:

- MOESI state machine for each cache line.
- Snoop response aggregation and forwarding.
- Error detection for invalid state transitions.

Key Code:

- State transition logic for MOESI protocol.
- Logic for aggregating and forwarding snoop responses.

nebula_niu_axi.sv: AXI NIU (1/2)

Design Decisions:

- Buffers and arbitrates between protocol and NoC domains.
- Handles packetization and depacketization.
- Manages local error and status reporting.

Key Code:

- Logic for buffering and arbitration between AXI and NoC.
- Packetization and depacketization logic.

Cool Features:

- Local error detection and reporting.
- Arbitration between multiple protocol sources.
- Parameterized for different AXI widths.

nebula_packet_assembler.sv: Packet Assembly (1/2)

Design Decisions:

- Converts protocol transactions to NoC flit packets.
- Segments payloads across multiple flits.
- Generates flit headers with CRC and sequence info.

Key Code:

- Header construction logic for different packet types.
- Payload segmentation and CRC generation logic.

nebula_packet_assembler.sv: Packet Assembly (2/2)

Cool Features:

- Address-to-coordinate mapping for mesh routing.
- CRC generation for each packet.
- Handles both single-flit and multi-flit packets.

nebula_packet_disassembler.sv: Packet Disassembly (1/2)

Design Decisions:

- Reconstructs protocol transactions from flits.
- Verifies CRC and sequence numbers.
- Handles out-of-order and multi-path delivery.

Key Code:

- Logic for reassembling transactions from received flits.
- CRC verification and error handling logic.

nebula_packet_disassembler.sv: Packet Disassembly (2/2)

Cool Features:

- Out-of-order packet reassembly.
- Error flagging for CRC mismatches.
- Sequence number tracking for reliability.

nebula_pkg.sv: Global Package (1/2)

Design Decisions:

- Centralizes all parameter, type, and enum definitions.
- Ensures consistency across all modules.
- Provides utility functions for coordinate mapping, CRC, etc.

Key Code:

- Parameter and type definitions for NoC and protocol.
- Utility function implementations (e.g., CRC calculation).

nebula_pkg.sv: Global Package (2/2)

Cool Features:

- Parameterized mesh and flit widths.
- Strong type safety for protocol and NoC structures.
- Utility functions for address/coordinate conversion.

nebula_top_simple.sv: Minimal Top (1/2)

Design Decisions:

- Minimal mesh and protocol integration for bring-up.
- Used for unit and smoke tests.

Key Code:

- Minimal mesh instantiation and AXI interface.
- Basic configuration and status logic.

nebula_top_simple.sv: Minimal Top (2/2)

Cool Features:

- Fast simulation for basic connectivity checks.
- Easy to extend for new protocol bridges.

common/: Reusable Modules (1/2)

Design Decisions:

- Implements core NoC utilities: CRC, FIFO, credit flow, arbitration.
- Parameterized for width, depth, and protocol.

Key Code:

- CRC calculation function and FIFO module.
- Credit flow control and round-robin arbiter.

common/: Reusable Modules (2/2)

Cool Features:

- Round-robin arbitration for fairness.
- Parameterized FIFO for flit/packet buffering.
- Modular, reusable code for all NoC blocks.

Key Code:

- Generic FIFO read/write logic.
- Arbiter grant and request handling.